



(a) textbook ABC-PMC, schematic:
one update per generation
 shown: generation $t \rightarrow$
 $t + 1$, after 135 simulations
 P_{t+1} : the N kept, weighted;
 the rejected are discarded



(b) generation-free (ours):
one update per arrival
 shown: the update after 84 evaluations
 θ^* to a free worker; ρ^* arrives later: append $(\theta^*, \rho^*, \epsilon_n, w^*)$



fewer than k particles: prior draw, $w^* = 1$, no bandwidth stored, straight to a worker