


(a) textbook ABC-PMC, schematic:

one update per generation

shown: generation $t \rightarrow t + 1$, after 135 simulations

P_{t+1} : the N kept, weighted;
the rejected are discarded




barrier: q_{t+1} needs the complete P_{t+1}

(b) generation-free (ours):

one update per arrival

shown: the update after 84 evaluations

θ^* to a free worker; ρ^* arrives
later: append $(\theta^*, \rho^*, \epsilon_n, w^*)$



fewer than k particles: prior
draw, $w^* = 1$, no bandwidth
stored, straight to a worker